

ORACLE

Oracle Database In-Memory

AskTOM Office Hours – 23c New Features – November 9, 2023

Andy Rivenes

Database In-Memory Product Manager

Twitter: @TheInMemoryGuy

Email: andy.rivenes@oracle.com



Database In-Memory 23c New Features

Oracle Database 23c Free

- All of the new features discussed today are available in Oracle Database 23c Free
- The 23c Free version is 23.3
- Examples were made using 23c Free in a Virtual Box environment
 - SGA_TARGET = 1400M
 - INMEMORY_SIZE = 800M (for AIM tests this was reduced to 400M)
- SSB schema from the Database In-Memory LiveLabs environment was used for most examples
- Oracle Database 23c Free is available here: <https://www.oracle.com/database/free/>
- Oracle Database 23c Free limitations:
 - 2 CPUs for foreground processes
 - 2GB of RAM (SGA and PGA combined)
 - 12GB of user data on disk (irrespective of compression factor)

23c Free Database In-Memory Features

- Changes to v\$inmemory_area
- New - Selective Columns
- Automatic In-Memory enhancements
- New – Automatic In-Memory Sizing
- In-Memory Deep Vectorization enhancements
- New – In-Memory Advisor
 - In-Memory Eligibility Tool
 - In-Memory Advisor (2.0)

Changes to v\$inmemory_area

- A new pool has been added to the In-Memory area
- This new pool takes most of the structures, and space required, out of the Shared Pool

Connected to:
Oracle Database 23c Free Release 23.0.0.0.0 - Develop, Learn, and Run for Free
Version 23.3.0.23.09

```
SQL> @09_im_usage.sql
Connected.
SQL> SELECT pool, alloc_bytes, used_bytes, populate_status, con_id
2 FROM v$inmemory_area;
```

POOL	ALLOC_BYTES	USED_BYTES	POPULATE_STATUS	CON_ID
1MB POOL	796,917,760	680,525,824	DONE	3
64KB POOL	25,165,824	4,128,768	DONE	3
IM POOL METADATA	16,777,216	16,777,216	DONE	3

SQL>



Selective In-Memory Columns

New INMEMORY Column Include/Exclude Syntax

- Now in 23c you can specify ALL columns as INMEMORY or NO INMEMORY and then selectively include or exclude columns from population:
 - ALTER TABLE <table name> NO INMEMORY(ALL) INMEMORY(col1, col2);
 - ALTER TABLE <table name> INMEMORY(ALL) NO INMEMORY(col1, col2);

Selective In-Memory Columns Example

```
SQL> alter table cust_test inmemory (c_custkey,c_name) no inmemory(all);
```

Table altered.

```
SQL> SELECT TABLE_NAME, COLUMN_NAME, INMEMORY_COMPRESSION from V$IM_COLUMN_LEVEL where TABLE_NAME = 'CUST_TEST';
```


TABLE_NAME	COLUMN_NAME	INMEMORY_COMPRESSION
-----	-----	-----
CUST_TEST	C_CUSTKEY	DEFAULT
CUST_TEST	C_NAME	DEFAULT
CUST_TEST	C_ADDRESS	NO INMEMORY
CUST_TEST	C_CITY	NO INMEMORY
CUST_TEST	C_NATION	NO INMEMORY
CUST_TEST	C_REGION	NO INMEMORY
CUST_TEST	C_PHONE	NO INMEMORY
CUST_TEST	C_MKTSEGMENT	NO INMEMORY

8 rows selected.

```
SQL>
```

Automatic Enablement of In-Memory Features

- In Oracle Database 23c AIM has been enhanced to add the ability to automatically:
 - Enable Join Groups where beneficial
 - Enable In-Memory Optimized Arithmetic for beneficial number columns
 - Enable higher compression levels on specific columns
 - Unpopulate unused columns
- Reduces manual effort required to leverage key performance features
- Enhanced workload analysis to better account for mixed workload environments
 - DML overhead with In-Memory (for example, fetching invalid rows in column store from buffer cache) is factored into analysis
- No application changes required



Automatic In-Memory Features

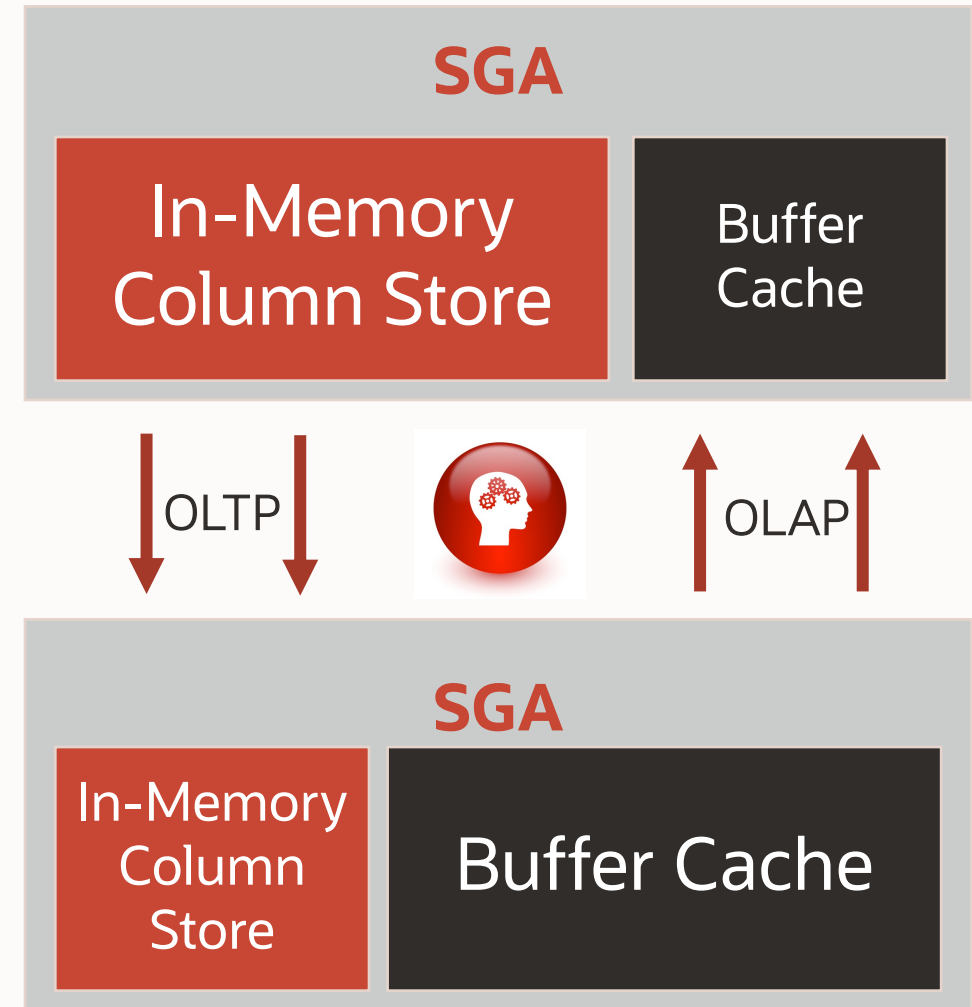
Control and Status

- DBMS_AUTOIM Package
 - SET_PARAMETER
 - Control feature creation (on or off)
 - Set statistics time window
 - ACTIVITY_REPORT
 - Generate AIM task activity for a specified time interval
- DBA_AIM_PERF_FEATURES View
 - Determine owner, table, columns that have AIM performance features enabled
 - Features:
 - Optimized Arithmetic
 - Bloom filter optimization (cached hash values)
 - Stored using vector optimization
 - Join Group

Automatic In-Memory Sizing

Automatic Sizing of the In-Memory Column Store

- Auto-size column store to factor in different workload types sharing resources at the same time
 - **Shrink** column store and increase buffer cache if DML intensive workload predicted.
 - **Grow** column store and shrink buffer cache if OLAP intensive workload predicted
- Works with Automatic In-Memory (AIM)
 - Requires AIM level HIGH and ASMM to be enabled (i.e. SGA_TARGET)
- The INMEMORY_SIZE parameter sets the minimum IM column store size
- Works in conjunction with In-Memory Hybrid Exadata Scans to ensure that partially populated objects can be accessed both in the IM column store and IM columnar format in Exadata Smart Flash Cache (i.e. cell memory)



Automatic In-Memory Sizing Example

Required parameter settings:

- heat_map = on
- sga_target
- inmemory_automatic_level = high

IM column store must be under memory pressure:

OWNER	SEGMENT_NAME	PARTITION_NAME	POPULATE_STATUS	Disk Size	In-Memory Size	Bytes Not Populated
SSB	CUSTOMER		COMPLETED	24,928,256	23,199,744	0
SSB	DATE_DIM		COMPLETED	122,880	1,179,648	0
SSB	LINEORDER	PART_1996	OUT OF MEMORY	565,018,624	263,061,504	267,354,112
SSB	LINEORDER		OUT OF MEMORY	563,470,336	210,501,632	323,739,648
SSB	LINEORDER		STARTED	563,322,880	131,072	563,322,880
SSB	PART		COMPLETED	56,893,440	16,973,824	0
SSB	SUPPLIER		COMPLETED	1,769,472	2,228,224	0



Automatic In-Memory Sizing Example, Part 2

```
select
  component, oper_type, oper_mode, parameter, initial_size,
  target_size, final_size,
  to_char(start_time,'MM/DD/YYYY HH24:MI:SS') start_time,
  to_char(end_time,'MM/DD/YYYY HH24:MI:SS') end_time
from
  v$sga_resize_ops
where
  component in ('shared_pool','DEFAULT buffer cache','In-Memory Area')
order by
  start_time
/
```

COMPONENT	OPER_TYPE	OPER_MODE	PARAMETER	INITIAL_SIZE	TARGET_SIZE	FINAL_SIZE	START_TIME	END_TIME
In-Memory Area	STATIC		inmemory_size	0	419430400	419430400	11/06/2023 18:11:09	11/06/2023 18:11:09
DEFAULT buffer cache	INITIALIZING		db_cache_size	226492416	226492416	226492416	11/06/2023 18:11:09	11/06/2023 18:11:09
DEFAULT buffer cache	STATIC		db_cache_size	0	226492416	226492416	11/06/2023 18:11:09	11/06/2023 18:11:09
In-Memory Area	GROW	DEFERRED	inmemory_size	419430400	545259520	545259520	11/06/2023 18:25:32	11/06/2023 18:25:32
DEFAULT buffer cache	SHRINK	DEFERRED	db_cache_size	226492416	100663296	100663296	11/06/2023 18:25:32	11/06/2023 18:25:32

In-Memory Vectorization Framework: Multi-Level Joins and Aggregations

- In-Memory Deep Vectorization Framework was introduced on Oracle Database 21c
 - The first feature that leveraged the framework improved the performance of single level hash joins
- In Oracle Database 23c the In-Memory Deep Vectorization Framework adds more join support:
 - Multiple Join Keys
 - Semi and Outer Join
 - Fully Grouping and Aggregation
 - Multiple Levels of Joins
- This translates to faster join performance by utilizing SIMD optimizations during join processing

Semi-Join Example

```
select /*+ MONITOR */ count(l.lo_custkey)
from lineorder l
where l.lo_partkey IN (select p.p_partkey from part p)
and l.lo_quantity <= 3;
```

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT					5113 (100)			
1	SORT AGGREGATE		1	18					
* 2	HASH JOIN RIGHT SEMI		1640K	28M	12M	5113 (8)	00:00:01		
3	TABLE ACCESS INMEMORY FULL	PART	800K	3906K		73 (3)	00:00:01		
4	PARTITION RANGE ALL		1640K	20M		2448 (15)	00:00:01	1	3
* 5	TABLE ACCESS INMEMORY FULL	LINEORDER	1640K	20M		2448 (15)	00:00:01	1	3



Semi-Join Example, Part 2

SQL*Plus way to determine Deep Vector usage (similar to Join Groups)

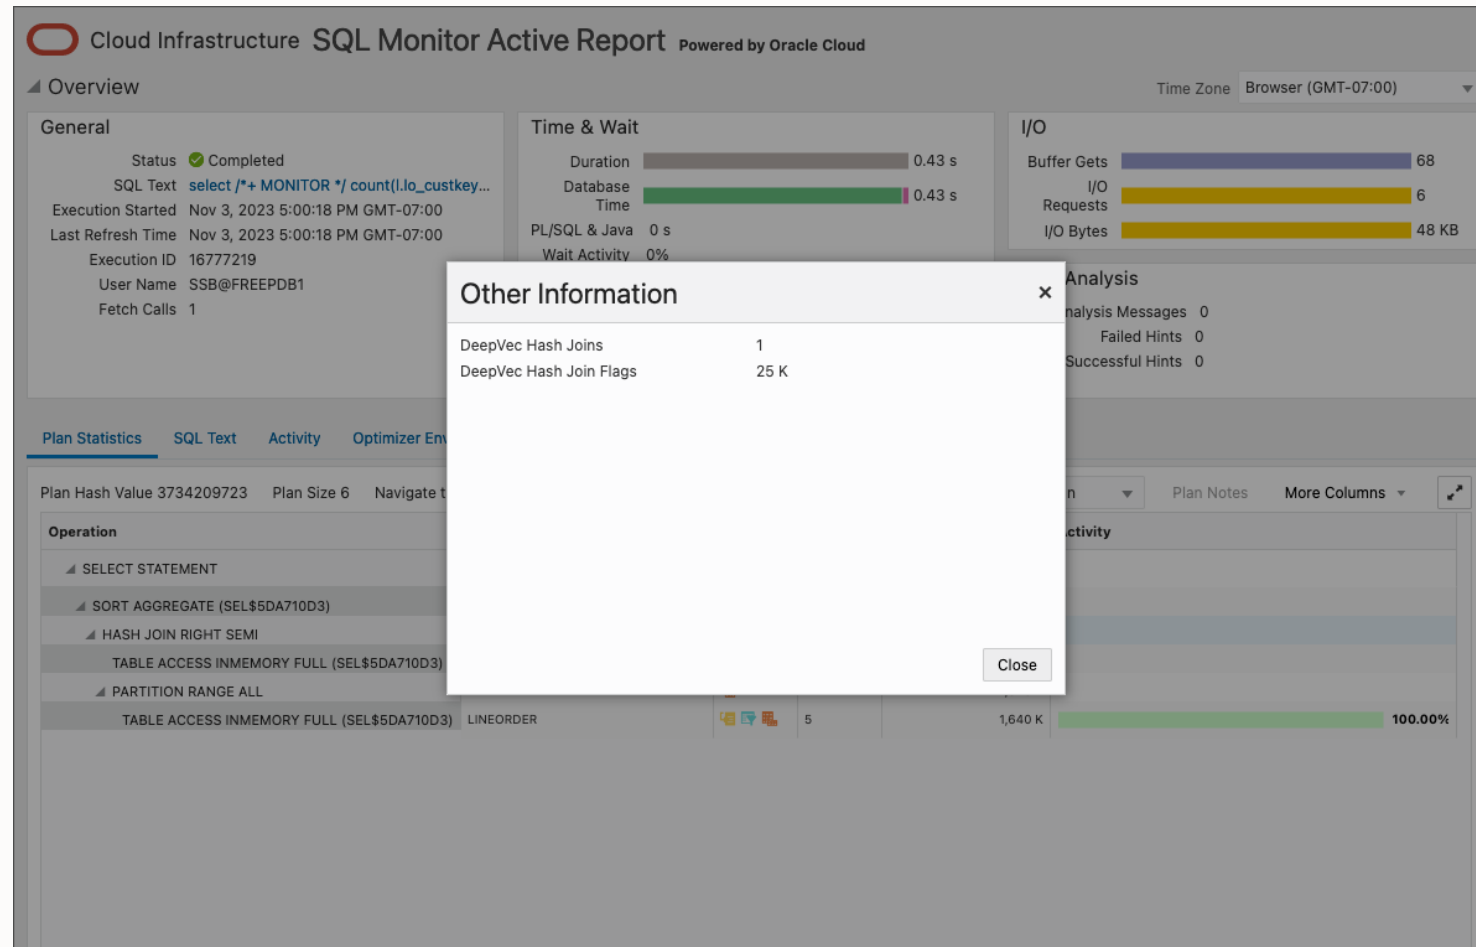
```
SELECT
  ' ' || deepvec.rowsource_id || ' - ' row_source_id,
  CASE
    WHEN deepvec.deepvec_hj IS NOT NULL
    THEN
      'deep vector hash joins used: ' || deepvec.deepvec_hj || ', deep vector hash join flags: ' || deepvec.deepvec_hj_flags
    ELSE
      'deep vector HJ was NOT leveraged'
  END deep_vector_hash_join_usage_info
FROM
  (SELECT EXTRACT(DBMS_SQL_MONITOR.REPORT_SQL_MONITOR_XML, q'#/operation[@name='HASH JOIN' and @parent_id]#') xmldata
   FROM   DUAL) hj_operation_data,
  XMLTABLE('/operation'
    PASSING hj_operation_data.xmldata
    COLUMNS
      "ROWSOURCE_ID"      VARCHAR2(5) PATH '@id',
      "DEEPEVEC_HJ"        VARCHAR2(5) PATH 'rwsstats/stat[@id="11"]',
      "DEEPEVEC_HJ_FLAGS"  VARCHAR2(5) PATH 'rwsstats/stat[@id="12"]') deepvec;
```

Deep Vectorization Usage:

2 - deep vector hash joins used: 1, deep vector hash join flags: 24576

Semi-Join Example, Part 3

SQL Monitor way to determine Deep Vector usage



Semi-Join Example, Part 4

Execution Plan Output

Plan hash value: 3734209723

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT					5113 (100)			
1	SORT AGGREGATE		1	18					
* 2	HASH JOIN RIGHT SEMI		1640K	28M	12M	5113 (8)	00:00:01		
3	TABLE ACCESS INMEMORY FULL	PART	800K	3906K		73 (3)	00:00:01		
4	PARTITION RANGE ALL		1640K	20M		2448 (15)	00:00:01	1	3
* 5	TABLE ACCESS INMEMORY FULL	LINEORDER	1640K	20M		2448 (15)	00:00:01	1	3

Predicate Information (identified by operation id):

2 - access("L"."LO_PARTKEY"="P"."P_PARTKEY")
5 - inmemory("L"."LO_QUANTITY"<=3)
filter("L"."LO_QUANTITY"<=3)

25 rows selected.

SQL>

SQL> set echo off

Deep Vectorization Usage:

2 - deep vector hash joins used: 1, deep vector hash join flags: 24576



New Database Embedded In-Memory Advisor

- Previous version
 - Standalone package that had to be installed in the database
 - Analyzed existing database workload using AWR/ASH data to determine benefit that Database In-Memory might provide
 - Generated an HTML report with estimates for overall benefit, benefit for top SQL statements and objects based on an in-memory size
- New in 23c, the In-Memory Advisor is now part of Oracle Database
- Relies on Heat Map data for analysis
 - Heat Map is now available as part of Oracle Database Enterprise Edition (no separate license)
- A new Eligibility Tool has been added to quickly identify databases where Database In-Memory would NOT be useful – Backported to 19c (19.20)
- A comprehensive advisor analysis based on a workload timeframe and an object benefit analysis

In-Memory Advisor Example

Eligibility Tool

```
SQL> variable inmem_eligible BOOLEAN
```

```
SQL> variable analysis_summary VARCHAR2(4000)
```

```
SQL> exec dbms_inmemory_advice.is_inmemory_eligible(26, 30, :inmem_eligible, :analysis_summary);
```

PL/SQL procedure successfully completed.

```
SQL> print inmem_eligible
```

```
INMEM_ELIGI
```

```
-----
```

```
TRUE
```

```
SQL> print analysis_summary
```

```
ANALYSIS_SUMMARY
```

```
-----
```

```
Observed Analytic Workload Percentage is 100% is greater than target Analytic Workload Percentage 20%
```

```
SQL>
```

In-Memory Advisor Example, Part 1

New Native In-Memory Advisor

```
SQL> variable taskid NUMBER;  
SQL> exec dbms_inmemory_advice.start_tracking(:taskid);
```

PL/SQL procedure successfully completed.

```
SQL> print taskid
```

```
      TASKID  
-----  
          1
```

```
SQL>
```

```
*** Do analytic work here ***
```

```
SQL> exec dbms_inmemory_advice.stop_tracking;
```

PL/SQL procedure successfully completed.

```
SQL>
```

In-Memory Advisor Example, Part 2

New Native In-Memory Advisor

```
SQL> exec dbms_inmemory_advise.generate_advise;
```

PL/SQL procedure successfully completed.

```
SQL> select round(INMEMORY_SIZE/1024/1024,3) as inmemory_size_mb,  
       trunc((ESTIMATED_DB_TIME_HIGH/60),2) as estimated_db_time_minutes,  
       recommended_obj_list as recommended_objects  
from dba_inmemory_advisor_recommendation  
where task_id = 4;
```

2	3	4	5
INMEMORY_SIZE_MB	ESTIMATED_DB_TIME_MINUTES	RECOMMENDED_OBJECTS	
0	10.18		
302	6.73	Owner: SCOTT Table: ANA_1 ; Owner: SCOTT Table: ANA_5 ; Owner: SCOTT Table: ANA_4 ;	
528	5.06	Owner: SCOTT Table: ANA_1 ; Owner: SCOTT Table: ANA_5 ; Owner: SCOTT Table: ANA_4 ; Owner: SCOTT Table: ANA_3 ;	
754	3.4	Owner: SCOTT Table: ANA_1 ; Owner: SCOTT Table: ANA_5 ; Owner: SCOTT Table: ANA_4 ; Owner: SCOTT Table: ANA_3 ; Owner: SCOTT Table: ANA_2 ;	



Where Can You Get More Information?

<https://blogs.oracle.com/in-memory/dbim-resources>

